

Programming Languages, Part A

▼ Status

Completed

▼ Course Motivation

01/14/24

Introduction

- On a theoretical level, all programming languages are equivalent in the sense that any concept can be implemented in any language.
 - Nonetheless, certain concepts can be taught easier in certain programming languages.
 - The universality of a concept can be emphasized by implementing that concept in multiple programming languages.
- A programming language is just an interface for writing software.
 - Therefore there is no substitute for precisely understanding programming language semantics.

Why study general programming language concepts?

- To understand programming languages, one needs to understand semantics and idioms.
 - Correct reasoning about programs, interfaces, and compilers requires knowledge of semantics.
 - Idioms make you a better programmer.
 - Seeing idioms in multiple settings clue you in to common programming patterns that can be implemented in most, if not all, languages.

Are all programming languages the same?

- On a technical level, all programming languages are the same.
 - Any input-output behavior implementable in language X is implementable in language Y (Church-Turing thesis).
- Some fundamentals reappear, such as variables, abstraction, one-of types, recursive definitions, and so on.
- What separates these programming languages are:
 - The fact that the primitive/default condition in one language is awkward in another.
- Learning other programming languages is exciting, and can make you appreciate your preferred programming languages even more.

Why functional languages?

- Functional languages are useful because:
 1. They have language features that are invaluable for correct, elegant, and efficient software, and are a great way to think about computation.
 2. They have always been “ahead of their time”.
 - a. Features like garbage collection, higher-order functions, type inference, and recursion were all implemented in functional languages, and dismissed as worthless ideas until they weren’t.

- b. Functional language concepts tend to get assimilated into other programming languages over time.
- 3. They are more well-suited towards where computing trends are going.
- With the absence of state, every input will have a consistent output.
 - Avoiding mutation is the easiest way to make concurrent and parallel programming easier.

Why SML, Racket, and Ruby?

- These three programming languages provide a very useful combination.

	dynamically typed	statically typed
functional	Racket	SML
object-oriented	Ruby	Java/C#/Scala

- SML provides polymorphic types, pattern-matching, and abstract types and modules.
- Racket provides dynamic typing, macros, and minimalist syntax.
 - We can define our own language in the Racket system.
- Ruby provides classes but not types, and is very object oriented.
- There are reasonable questions to ask when deciding to use/learn a programming language:
 - What libraries are available?
 - What tools are available?
 - What can land a job?
 - What is the industry standard?
- Not needing to deal with most of these questions for the course affords us the opportunity to use these more niche programming languages to learn the fundamentals.
 - In the long term, we will be able to apply these fundamentals to any programming language we want.

▼ Lecture 1:

01/05/24

ML variable bindings and expressions

- Comments in SML are created by enclosing text with `(* *)`.

```
(* This is a comment *)
```

- Variables are created using the `val` keyword.
 - Semicolons `;` are used to mark the end of expressions.

```
(* val <x> = <expression> *)

(* Assign 34 to x *)
val x = 34;
```

- Files can be run in the SML (read-eval-print-loop) REPL with `use`.

```
(* Run fil.sml *)

$ use "file.sml";
```

- Type-checking** means checking what types are allowed for expression, and what type expression return.
 - If a type is invalid, an error will be thrown.
 - SML is a type-checked language.
- The **static environment** is active before a program is run.
 - The static environment does type-checking among other tasks.
- The **dynamic environment** is active when a program is run.
 - The dynamic environment evaluates bindings and expressions.
- New environments are “created” as the program is evaluated.
- Conditionals can be created using the keywords `if`, `then`, and `else`.

```
(* Evaluate the absolute value of z *)

val abs_oz_z = if z < 0 then 0 - z else z;
```

- Syntax** refers to how something is written, where **semantics** refers to what something means.
- In this case, semantics are broken down into:
 - Type checking (before the program runs).
 - Evaluation (as the program runs).

Rules for expressions

- For every expression, ask yourself three questions:
 - What is the **syntax**?
 - What are the **type-checking rules**?
 - How could the type-checking fail?
 - What are the **evaluation rules**?
 - An expression produces a **value**.
 - All values are expressions, but not all expressions are values.

The REPL and errors

- We need to ask ourselves some questions:
 - How do we run programs using the REPL?
 - What happens when we make mistakes?
- `use` enters variable bindings from a `.sml` file.
 1. The REPL `reads` the program.
 2. Next it `evaluates` the program.
 3. Finally, it `prints` the results.
 4. Repeat
- Values are printed in the REPL with their name, value, and type.

```
(* Anatomy of a value in the SML REPL *)

val <name> = <value> : <type>

(* Example of "x" being bound to the integer 7 *)

val x = 7 : int
```

- Errors can be raised as a result of improper syntax, type-checking, or during evaluation.
- The `~` character is used in SML to negate numbers.

```
(* negative 5 *)

~5
```

- The keyword `div` is used in SML to divide numbers.

```
(* 1 divided by 2 *)

1 div 2
```

Shadowing

- **Shadowing** is when a variable is added to an environment more than once.
 - In SML, assignment to a variable is not possible (they are immutable).
 - One can only shadow it in a later environment.

```
(* "a" is being shadowed, as its change in value doesn't affect other
   assignment statements like with "b" *)

(* In this example, the previous "a" is "shadowed" by the most recent "a" *)
```

```

val a = 10

(* a : int
   a -> 10 *)

val b = a * 2

(* a -> 10
   b -> 20 *)

val a = 5 (* this is not an assignment statement *)

(* a -> 5, b -> 20 *)

```

- Expressions in variable bindings are evaluated **eagerly**, meaning it is evaluated before the variable binding finishes.
 - After this, the expression producing the value is irrelevant.

Functions

- Functions are a key building block in SML, just like most other programming languages.

```

(* Example function to raise "x" to the power of "y" *)

(* Anatomy of a function *)
fun <name>(<param> : <type>...) =
  <expression>

(* Call a function *)
<name>(<arg>...)

fun pow(x : int, y : int) =
  if y = 0
  then 1
  else x * pow(x, y - 1)

```

- Functions are printed in the REPL with a special value `fn`, along with their name, input and output types.
 - Each input type is separated by a `*`.

```

(* Anatomy of a function in the REPL *)
val <name> = fn : type_1 * type_2 * ... * type_n -> type

(* Example using the previous "pow" function *)
val pow = fn : int * int -> int

```

- Remember that a function is a value itself.

Pairs and other tuples

- **Tuples** are a fixed number of values that may have different types.
- **Pairs** are tuples with two values.
- The type for tuples are their base component values separated by `*` s.
- The tuple components can be accessed individually using the syntax `#n`.

```
(* Example tuple *)
(* Tuple type is: int * bool *)

tup = (100, true)

(* Access the first element, 100 *)
#1 tup

(* Access the second element, true *)
#2 tup
```

Lists

- Unlike tuples, lists can have any number of elements.
- Unlike tuples, all element in a list have to contain the same type.

```
(* Empty list *)
[]

(* Anatomy of a list *)
[<val_1>, <val_2>, ..., <val_n>] (* where each value is of the same type *)
```

- `::` can be used to append a value to the front of a copy of a list.

```
(* Append "e1" onto "e2", creating a new list with "e1" on the front *)
e1::e2
```

- `null` returns a true if a list is empty.
- `hd` gets the first element of a list, or the head.
- `tl` get every element but the first element of a list, or the tail.

```
(* Ask if a list "e" is empty *)
null e

(* Produce "v1" of list "e" *)
hd e
```

```
(* Produce "v2, ..., vn" of list "e" *)
tl e
```

- The type `list` describes a list

```
(* A list "x" that bound to a list of ints, with type "int list" *)
$ val x = [1, 2, 3, 4] : int list
```

- An empty list has type `'a list`, or alpha list.
 - This means that the `'a` “alpha” can be replaced with any value.

Let expressions

- **Let expressions** allow for **local scoping** in SML.
- Let expressions are constructed in SML using the `let`, `in`, and `end` keywords.
 - The type returned by the expression is the type returned by the `in` binding enclosed the let expression.

```
(* Anatomy of a let expression *)
let
  <b_1,
    b_2,
    ...,
    b_n>
in
  <expression>
end
```

Nested functions

- Nested helper functions are useful when:
 - The helper function is unlikely to be useful elsewhere.
 - The helper function is likely to be misused elsewhere.
 - The helper function is going to be changed or removed later.
- One should but a function is a wide enough scope to save effort and avoids bugs, but also put it in a narrow enough scope to restrict how it is used.

01/06/24

Let and efficiency

- Avoid repeated recursive calls in a function, as this leads to exponential time complexity.
- To avoid this, `let` expressions can be used to precompute recursive function values.

```
(* "max_tail" is precomputed inside the let statement, and therefore the
recursive function is only called once per recursion level instead of
twice *)
```

```

(* Find the maximum value in a list *)
fun max (xs : int list) =
  if null xs
  then 0 (* horrible style, fix later in options section *)
  else
    let
      val max_tail = max(tl xs)
    in
      else if null (tl xs)
      then hd xs
      else if hd xs > max_tail
      then hd xs
      else
        max_tail
    end
end

```

Options

- Sometimes an function runs into a problem where it needs to return either a zero or one element list.
- The type `option` achieves this goal by giving “options” to how many values it can hold.
- `NONE` builds an empty option.
- `SOME` builds a one element option.
- `isSome` returns true if the option is a `SOME`.
- `valOf` returns the value of a `SOME`.

```

(* The type "int option" is returned, meaning either an "int" or "NONE"
   is returned *)

```

```

(* fn : int list -> int option *)
fun max (xs : int list) =
  if null xs
  then NONE
  else
    let
      val tl_ans = max1(tl xs)
    in
      if isSome tl_ans andalso valOf tl_ans > hd xs
      then tl_ans
      else
        SOME (hd xs)
    end
end

```


Booleans and comparison operations

- Boolean operations in SML are `andalso`, `orelse`, and `not`.
- For evaluating `e1 <boolean operation> e2`:
 - `e2` is not evaluated if `andalso e1` is `false`.
 - `e2` is not evaluated if `orelse e1` is `true`.
 - This means that `andalso` and `orelse` are keywords, where `not` is just a function.
- This means that program languages don't actually need Boolean keywords.
 - Nonetheless, these Boolean operations are useful as they are more concise and readable.

```
(* Equivalent to "andalso", "orelse", "not" *)
```

```
(* e1 andalso e2 *)
if e1
  then e2
else
  false
```

```
(* e1 orelse e2 *)
if e1
  then true
else
  e2
```

```
(* not e1 *)
if e1
  false
else
  true
```

- In SML, “not equal to” is `<>`.
- `<`, `>`, `<=`, and `>=` can not be used with a combination of `int` and `real`.
 - `real` is the SML type for floats.

Benefits of no mutation

- In SML, there is no mutation for most forms of data.
- A lack of mutation in SML means that **no data can be changed in a new dynamic environment**.

```
(* In SML, there is no way to track if the returned "pr" is an alias
  of the input "pr" or not, where it would be in a language with
  mutable tuples *)
```

```
fun sort_pair (pr : int * int) =  
  if #1 pr < #2 pr  
  then pr  
  else  
    (#2 pr, #1 pr)
```

- Without mutation:
 - No code can ever distinguish aliasing vs. identical copies.
 - Because the original data can never change, there is no need to think about aliasing.
 - We can use aliasing to save space, without danger.
- This leads to SML being more efficient than other programming languages in some cases.
 - For example, `tl` in SML is constant in time complexity since it doesn't need to copy the list every time a recursive function is called.

Pieces of a programming language

1. **Syntax**: How do you write language constructs?
 2. **Semantics (Evaluation rules)**: What do programs mean?
 3. **Idioms**: What are typical patterns for using language features to express your computation?
 4. **Libraries**: What facilities does the language provide “standard”?
 5. **Tools**: What do language implementation provide to make your job easier?
 - a. These tools are separate from the language itself.
- In practice, all of these are essential for good programmers.
 - It is essential how each of these pieces are separate from one another.
 - This course **focusses on semantics and idioms**.
 - Learning semantics and idioms can make you better at understanding libraries and tools.

▼ Lecture 2:

01/07/24

Building compound types

- Most programming languages have **base** and **compound** types.
 - There are three important type building blocks that can describe most types:
 - **Each-of** types, **one-of** types, and **self-reference** types.
 - Tuples build each-of types.
 - Options build one-of types.
 - Lists use all three building blocks.
 - These types can be nested to create more complex types.
-

Records

- **Records** are a type that can have multiple field names, with each field having a unique type.
 - Records are enclosed by `{ }` with field names separated by commas.
 - Field names can be accessed with `#`.

```
(* A record "x" with the type {foo=int*int, bar=bool} *)
val x = {foo=(3, 2), bar=true}

(* Bind (3, 2) to y *)
val y = #foo x
```

- Tuples are useful when one wants to refer to elements by position, where records are useful when one wants to refer to elements by field name.\

Tuples as syntactic sugar

- There are no such thing as tuples in SML, as they are actually just syntactic sugar for records with ordered numeric field names.

```
(* This record will be show as a tuple (4, true, 7) with type int * bool * int
   by the REPL *)
val actually_a_record = (1=4, 3=7, 2=true)

(* 4 *)
#1 actually_a_record
```

- **Syntactic sugar** simplify the understanding and implementation of a programming language.
- It also makes the core of programming languages more concise, and shortens proofs.
 - It is only “syntactic”, because the semantics can still be described in terms of another construct.
 - It is “sugar” because it makes the language sweeter 🍬.
 - As discussed previously, `andalso` and `orelse` are also forms of syntactic sugar.

01/09/24

Datatype bindings

- A **datatype** is a one-of type in SML.
- **datatype** s add **constructors** to the environment.
 - One job of a constructor is to be a function that makes values of the new type.

```
(* mytype contains one of int * int, string, or mytype, with each one
   being a constructor function *)
(* Each is a function that takes a value of the specified type, and converts
   it to a "mytype" *)
```

```
(* TwoInts: int * int -> mytype *)
(* Str: string -> mytype *)
(* Pizza: mytype *)

datatype mytype = TwoInts of int * int | Str of string | Pizza
```

- The value contains a **tag** and its corresponding **data**.
 - The tag specifies which constructor is being used.
- There are two aspects to accessing a datatype value.
 - Check what variant it is (what constructor made it).
 - Extract the data (if that variant has any).
- Types like lists and options have these components:
 - `null` and `isSome` check variants.
 - `hd`, `tl`, and `valOf` extract data.

Case expressions

- **Case expressions** are used to extract the data from `datatype` bindings.
 - It does so through accessing the pieces of a `datatype`.
 - It is really a multi-branch conditional to pick a branch based on a variant.
 - Case expressions use the first branch that matches via **pattern matching**.
 - They extract data and binds that data to local variables in the branch.
 - All branches must have the same type.

```
(* A case expression takes a value "x" with the type "datatype", and
   access the pieces of "mytype". *)

(* mytype -> int *)
fun f (x : mytype) =
  case x of
    Pizza => 3
  | Str s => s ^ "s"
  | TwoInts (i1, i2) => i1 + i2

(* Call "f" with a "mytype", return 3 *)
f (TwoInts (1, 2))
```

- A **pattern** in SML is a constructor name followed by the right number of variables.
 - Syntactically, most patterns **look like expressions, but they are not the same**.
 - They are not evaluated, they are just used for pattern matching and accessing `datatype` variables.
 - Patterns are evaluated in order.

Useful datatypes

- `datatype` s have many applications, such as enumerations, or representing real-world objects.

```
(* An enumeration to describe playing card suits *)
datatype suit = Club | Diamond | Heart | Spade

(* Datatype to describe students in a student database *)
datatype id = StudentNum of int
           | Name of string * (string option) * string
```

- `datatype` s are different from records in the sense that they are one-of types instead of each-of types.
- `datatype` s can have recursive constructors.

```
(* The "exp" datatype uses itself in the constructor of "Add" and "Negate" *)

datatype exp = Constant of int
           | Negate    of exp
           | Add      of exp * exp
```

Type synonyms

- `datatype` bindings introduce a new type name that is distinct from all existing types.
- The only way to create values of the new type is a constructor.
- A `type synonym` is a new kind of binding, which creates another name for an existing type.
 - The type and name are interchangeable in every way.

```
(* The type "card" is a synonym for a tuple of type "suit" and type "int" *)

type card = suit * int

(* Create a new card value representing Ace of spades *)
(* This value is both a "card" and a "suit * int" *)

val ace_of_spades : card = (Ace, 1)
```

- For now, type synonyms are just a convenience for talking about types.

Lists and options are datatypes

- Options and lists are just a predefined datatype binding.
 - For options, `NONE` and `SOME` are constructors, and `isSome` and `valOf` use pattern matching.
 - For lists, `[]` and `::` are constructors.
- Pattern matching is better for options and lists for the same reasons as for all datatypes.

- There are no missing cases.
- There are no exceptions for the wrong variant.
- `null`, `tl`, `[]`, etc. are predefined for passing them as arguments to other functions, or just plain convenience.

Polymorphic datatypes

- **Polymorphism** is the ability for a value to take on multiple forms.
- Datatype bindings like `list` and `option` are polymorphic type constructors, in the sense that `list` is not a type, but type `'a list` is a type.
- SML allows one to define their own polymorphic type constructors.
- `datatype` has the option to take one or more type parameters in constructor definitions.

```
(* Example of a custom option datatype that takes in type 'a *)

datatype 'a option = NONE | SOME of 'a

(* Example of a tree datatype that takes in types 'a and 'b *)

datatype ('a, 'b) tree = Node of 'a * ('a, 'b) tree * ('a, 'b) tree
                      | Leaf of 'b
```

Each-of pattern matching

- All `val` and `fun` binding uses pattern matching.
- Pattern matching works for records and tuples, as in the pattern `(x1, ..., xn)` matches the tuple `(v1, ..., vn)`, and the pattern `{f1=x1, ..., fn=xn}` matches the pattern `{f1=v1, ..., fn=vn}`.

```
(* Example of a case expression to match a triple datatype and sum
   it three values *)

fun sum_triple triple =
  case triple of
    (x, y, z) + x + y + z
```

- Value bindings are actually `val p = e`, where `p` can be a pattern, not just a variable.

```
(* Building on the previous example, we now use values to extract the
   values from triple *)

fun sum_triple triple =
  let
    (* The type checker would throw an error if this pattern did not match *)
    val (x, y, z) = triple
  in
```

```
x + y + z
end
```

- A function argument can also be a pattern, meaning that the arguments `p` in a function call `f` can be matched with `fun f (p) = e .`

```
(* Building on the previous example, we now see that parameters in
   functions are just syntactic sugar that use pattern matching *)
(* Because a pattern is being used, not explicit types
   are needed for functions *)
```

```
fun sum_triple (x, y, z) =
    x + y + z
```

- This removes the need for explicit types to be stated for arguments, as the type checker will catch any patterns that don't match their specified types.
- The important thing to understand is that this is exactly the same syntax as functions describes earlier in the course.
- We have been using pattern matching for functions all along.
 - Every function takes **exactly** one argument, that argument being a pattern.
 - What we have been calling multi-argument functions are just functions with a single tuple, and then pattern matching is used on that tuple to extract its values.
- This enables us to do useful things in SML that one may not be able to do in other programming languages.

```
(* Simple function to rotate the pieces of a tuple *)

fun rotate_left (x, y, z) = (y, z, x);

(* Because every function only takes one parameter, the return value can be
   immediately passed to another function, as long as the types match *)

rotate_left(rotate_left(3, 4, 5))
```

01/10/24

Type inference

- When case expressions are used for pattern matching, no types need to be explicitly stated in function parameters.

```
(* Using pattern matching is superior because the type checker knows how
   many fields the tuple has *)
```

```
fun sum_triple1 (x, y, z) =
    x + y + z
```

```
(* Without the int * int * int value, the type checker has
   no idea how many fields triple has *)
```

```
fun sum_triple2 triple =
  #1 triple + #2 triple + #3 triple
```

- If types are not used in a function and is flexible as a type, the type checker will declare it as an `'a` type.

```
(* Because "y" is not used, it does not necessarily have to be an int type,
   and therefore the type checker declares it as an 'a type *)
```

```
fun partial_sum (x, y, z) =
  x + z
```

Polymorphic and equality types

- `'a` is more general type that can be used in place of a more specific type, as long as the alpha is used consistently for only that specific type.
- `'a` in a function signature cannot be replaced with multiple types.
- Variables such as `''a` with two `'` specify equality types.
 - **Equality types** are types that you can use the `=` operator on.
 - For example, `fn` or `real` types do not count as equality types.
 - Equality types are more specific than polymorphic types.

```
(* The type signature for "same_thing" is ''a * ''a -> string *)
fun same_thing(x, y) =
  if x=y then "yes" else "no"
```

Nested patterns

- Patterns can be put inside of other patterns.
- Nested pattern can lead to elegant, concise code.
 - One should avoid nested case expressions when possible.
 - A common idiom is matching against a tuple of datatypes to compare them.
- `_` is a pattern “wildcard” that matches everything.
 - Wildcards are good style, as they specify that the data from a variable does not need to be used inside that branch.

```
(* Each list is matched with a head and a tail *)
(* The heads are extracted and the tails are recursively fed into the zip
   function until only empty lists are left *)

fun zip list_triple =
  case list_triple of
    ([], [], []) = []
```



```

    | (hd1::tl1, hd2::tl2, hd3::tl3) => (hd1, hd2, hd3)::zip(tl1, tl2, tl3)
(* If the length of the functions don't match, that is, a call has neither
   a fully empty list or a full list, the _ (default case) will be called
   and raise an error to inform that the lengths of the lists mismatch *)
    | _ => raise ListLengthMismatch

```

Function pattern syntactic sugar

- SML allows functions that only contain case expressions to be written more concisely.

```

(* Without syntactic sugar *)
fun f x =
    case x of
        p1 => e1
      | p2 => e2
      | ...
      | pn = en

(* With syntactic sugar *)

fun f p1 = e1
  | f p2 = e2
  | ...
  | f pn = en

```

Exceptions

- **Exceptions** are thrown when there are runtime conditions that should be an error.
- The `exception` keyword can be used to bind custom exceptions.
 - All exceptions in SML have the type of `exn`.
 - These exceptions have the option to carry values.

```

(* This exception binding introduces a new kind of exception *)

exception CustomException

(* This exception binding carries two ints *)

exception TwoIntException of int * int

```

- The `raise` keyword is used to throw exceptions in SML.

```

(* Implementation of "hd" *)
(* An error is raised when the list is empty and there is no hd to access *)

fun hd xs =

```

```
case xs of
  [] => raise List.Empty
| x::_ => x
```

- `handle` can be used to match an exception name to a pattern, and then execute an expression in that branch.

```
(* If the exception is not raised,
   the first expression is evaluated normally *)
(* If the exception is raised,
   then the second expression is evaluated instead *)
```

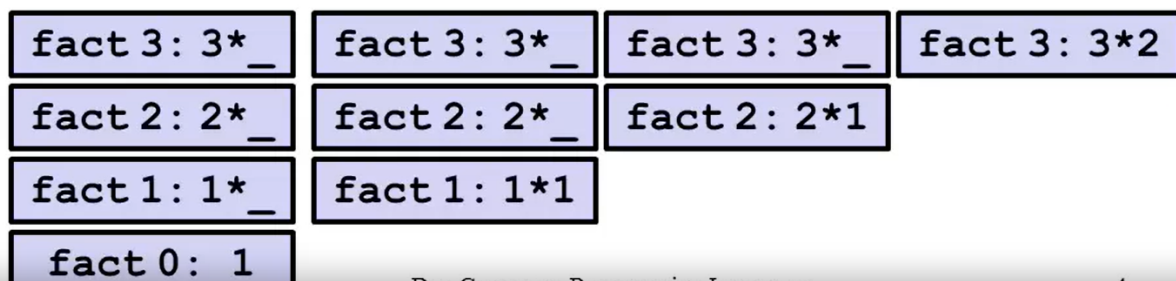

```
<expression1> handle
  <exception> => <expression2>
```

- Exceptions are very similar to `datatype` constructors.
 - Declaring an exception creates a constructor for type `exn`, which can be passed anywhere.

Tail recursion

- While a program runs, there is a `call stack` of functions that have started but not yet returned.
 - Calling a function `f` pushes an instance of `f` onto the stack.
 - When a call to `f` finishes, it is popped from the stack.
- These stack-frames store information like the value of local variable and “what is left to do” in the function.
- Due to recursion, multiple stack-frames may be calls to the same function.

```
fun fact n = if n=0 then 1 else n*fact(n-1)
val x = fact 3
```




- It is unnecessary to keep around a stack-frame just so it can get a callee's result and return it without any further evaluation.
- In that case, tail recursion can be implemented with the assistance of an accumulator.
 - Such a situation is called a **tail call**.
 - ML recognizes these tail calls and treats them differently to optimize them.
 - The caller is popped from the stack before the call, allowing the callee to reuse the same stack space.
 - With this, among other optimizations, recursion is as efficient as a loop.

```

fun fact n =
  let fun aux(n,acc) =
        if n=0
        then acc
        else aux(n-1,acc*n)
      in
        aux(n,1)
      end
  val x = fact 3

```



- It is reasonable to assume that all functional-language implementations do tail-call optimization.

Accumulators for tail recursion

- There is a methodology that can guide implementing an accumulator.
 1. Create a helper function that takes an **accumulator**.
 2. The old base case becomes the initial accumulator.
 3. The new base case becomes the final accumulator.
- Tail recursion can be implemented for case expressions.

```

(* The accumulator "accumulates" the sum of the int list *)

fun sum xs =
  let
    fun aux (xs, acc) =
      case xs of
        [] => acc
      | x::xs; => aux(xs', x + acc)
  in

```

```
    aux(xs, 0)
end
```

Perspective on tail recursion

- There are certain cases where tail-recursive functions cannot be evaluated in a constant amount of space.
 - For example, functions that process trees should utilize the call stack to keep track of where they have been in the tree.
- Beware of premature optimization.
 - One should favor clear, concise code.
 - Do use less space if inputs may be large.
- A tail call is a function call that is in **tail position**.
 - If an expression is not in tail position, then none of its subexpressions will be in tail position either.
 - The **tail position** is that last expression in a function that needs evaluating.

▼ Lecture 3:

01/12/24

First-class functions

- **Functional programming** means:
 1. Avoiding mutation in most/all cases.
 2. Using functions as values.
- This style encourages recursion and recursive data structures, and is closer to mathematical definitions.
- **First-class functions** can be used wherever we use other values.
 - Functions are values too.

```
(* Define some functions *)
fun double x = 2 * x
fun incr x = x + 1

(* Bind the functions to a tuple *)
(* val a_tuple = (fn, fn) : (int -> int) * (int -> int) *)
val a_tuple = (double, incr)

(* Access double and pass in 9 as an argument, results in 18 *)
val eighteen = (#1 a_tuple)(9)
```

- The most common use is as an argument or result of another function.
 - The function that takes or returns first-class functions is called a **higher-order function**.

- **Function closures** are functions that use bindings from outside the function definition, in the scope where the function is defined.

Functions as arguments

- We can pass one function as an argument to another function.
- First class functions allow us to create higher-order functions that generalize specific functions, and avoid repetition.

```
(* Example of an higher-order function that can do the tasks of multiple
   simple recursive functions *)

(* Apply a function to "x", "n" time *)
fun n_times(f, n, x) =
  if n=0
  then x
  else f(n_times(f, n-1, x))

(* Use the higher-order function to create a simple double "n" times function *)
fun double x = x * 2
fun double_n_times(n, x) =
  n_times(double, n, x)

(* Use the higher-order function to create a simple triple "n" times function *)
fun triple x = 3 * x
fun triple_n_times(n, x) =
  n_times(triple, n, x)
```

Polymorphic types and functions as arguments

- Higher-order functions are often so generic and reusable that they have polymorphic types.

```
(* The polymorphism of n_times makes it more useful *)
(* Type:
   ('a -> 'a) * int * 'a -> 'a *)

fun n_times(f, n, x) =
  if n=0
  then x
  else f(n_times(f, n-1, x))
```

- That being said, not all higher-order functions are polymorphic, and not all polymorphic functions are higher-order.

Anonymous functions

- **Anonymous functions**, commonly known as lambda functions in other programming languages, are functions that can be defined without a binding, but with an expression.
 - These are commonly used when a simple function is only needed in one place and nowhere else.

- Anonymous functions use `fn` instead of `fun`, use `=>` instead of `=`, and have no name.
- Anonymous functions **cannot define recursive functions**.
 - If our function doesn't have a name, therefore there is no way to invoke the function recursively.

```
(* Anonymous function anatomy:
   fn <args> => <expression>

(* We can define triple in triple_n_times with an anonymous function *)
fun triple_n_times(n, x) =
  n_times(fn x => 3 * x, n, x)
```

Unnecessary function wrapping

- Sometimes we don't need anonymous functions.
 - When all an anonymous function does is pass its argument to another function, just use that other function instead.

```
(* This lambda function is redundant *)
fun nth_tail(n, xs) =
  n_times(fn y => tl y, n, xs)

(* A simpler solution is just to use tl, as tl is
   already the expression that is needed *)
fun nth_tail(n, xs) =
  n_times(tl, n, xs)
```

Map and filter

- `map` and `filter` are higher-order functions that are built into most programming languages.
- `map` uses, or “maps” a function to every element in a list.
 - In SML, the function syntax is `List.map`.

```
(* How the built-in map function works *)
(* Type:
   (a' -> b') * 'a list -> 'b list *)
fun map (f, xs) =
  case xs of
    [] => []
  | x:xs' => (f x)::map(f, xs')

(* Use map, produce [5, 9, 13, 17] *)
(* Specific Type (with anonymous function):
   (int' -> int') * 'int list -> 'int list *)
val x1 = map(fn x => x + 1, [4, 8, 12, 16])
```

- `filter` filters a list based on whether each element returns true for the argument function.
 - In SML, the function syntax is `List.filter`.

```
(* How the built-in filter function works *)
(* Type:
   ('a -> bool) * 'a list -> 'a list *)
fun filter (f, xs) =
  case xs of
    [] => []
  | x:xs' => if (f x)
              then x::filter(f, xs')
              else filter(f, xs')

(* Specific Type (with anonymous function):
   ('int -> bool) * 'int list -> 'int list *)
(* Use filter, produce [12, 16] *)
val x2 = filter(fn x => x > 9, [4, 8, 12, 16])
```

01/14/24

Generalizing prior topics

- First class functions are also useful in the fact that they can return specific function based on input.

```
(* Returning a double or a triple function based on the input "f" *)
fun double_or_triple f =
  if f 7
  then fn x => 2 * x
  else fn x => 3 * x
```

- Higher order functions are not just good for numbers and lists.
 - They work great for common recursive traversals over data structures.
- Functions that processes data into Booleans are sometimes called **predicates**.

Lexical scope

- Function bodies can use any bindings that were defined **in their scope**, not the scope where the function is being called.
- The semantics for this is called **lexical scope**.
 - Understanding lexical scope is crucial for utilizing first-class functions effectively.

```
(* f uses the value of x = 1 that was mapped
   when the function was defined, not x = 2 *)

val x = 1
  (* x maps to 1 *)
fun f y = x + y
```

```

    (* f maps to a function that adds 1 to its argument *)
val x = 2
    (* x shadows to 2 *)
val y = 3
    (* y maps to 3 *)
val z = f (x + y)
    (* Call the function defined on line 2 with f(5) *)
    (* z maps to 6 *)

```

- The language implementation uses closures to hold on to older environments that wouldn't be around otherwise, in order to retrieve the correct function definition.
- A function value has two parts:
 1. The code.
 2. The environment that was current when the function was defined.
- These pieces of this pair cannot be accessed, but the pair can be called.
- This pair is called a **function closure**.
 - A call **evaluates the code using the environment**.

```

(* When we define "f", a function closure is created *)

(* Code: take y and have body x + y *)
(* Environment: x maps to 1 *)

val x = 1
fun f y = x + y
...

(* Now when "f" is used, it calls the closure where x is mapped to 1 *)

val x = 2
val z = f (x + y)

```

Lexical scope and higher-order functions

- The function closure binds values within a function's lexical scope, similar to how a **let** statement works.

```

(* Written code *)

val x = 1
fun f y = x + y

(* f is implemented with a let statement where x is defined within scope *)

fun f y =
  let

```



```
        x = 1
    in
        x + y
    end
```

Why lexical scope

- **Lexical scope** is the environment where a function is defined.
- **Dynamic scope** is the environment where a function is called.
- Lexical scope is used in almost all modern programming language implementation for three precise, technical reasons:
 1. Function meaning does not depend on variable names used.
 - a. Variable names can be changed within the scope, and it won't effect any of the function calls.
 - b. This allows for modularity in software, as variable names won't override one another if they are in separate lexical scopes.
 2. Functions can be type-checked where they are defined.
 - a. This can lead to operating on combinations of types that are not allowed, resulting in undefined types.
 3. Closures can easily store the data they need.
 - a. Variable names are "locked in" so-to-say.

```
(* A function that takes x, and returns a function that asks
   if y is greater than x *)
fun greater_than_x x = fn y => y > x

(* The value of x is locked in as -1 *)
(* fn y => y > -1 *)
greaterThan ~1
```

Closures and recomputation

- Closures can be used to avoid repeating computations **that do not depend on function arguments**.
 - If we did not use closures to store these computations, we would have to compute them every time the function is called instead.

```
(* Filter strings that are equal to or longer than string s *)
fun all_shorter_than_1 (xs, s) =
    filter (fn x => String.size x < String.size s, xs)

(* Avoid recomputation by computing the size of s once *)
fun all_shorter_than_2 (xs, s) =
    let
        val len_s = String.size s
    in
```

```

    filter (fn x => String.size x < len_s, xs)
end

```

Fold and more closures

- `fold` is an iterator over recursive functions.
 - It accumulates an answer by repeatedly applying `f` to an answer so far.

```

(* How the built-in fold function works *)
(* Type:
   ('a * 'b -> 'a) * 'a * 'b list -> 'a *)
fun map (f, acc, xs) =
  case xs of
    [] => acc
  | x:xs => fold(f, f(acc, x), xs)

(* Use fold to find the sum of a list, produce 10 *)
(* Specific Type:
   ('int * 'int -> 'int) * 'int * 'int list -> 'int *)
fold(fn (x, y) => x + y, 0, [1, 2, 3, 4])

```

- This pattern separates recursive traversal from data processing.
 - One can use the same traversal for different data processing.
 - We can reuse the same data processing for different data structures.
 - In both cases, *using a common vocabulary concisely communicates intent*.
- Functions like `map`, `filter`, and `fold` are much more powerful thanks to closures and lexical scope.
 - Functions passed to iterators can use any “private” data in its environment.
 - The iterator does not need to know what data is there or what type the data has.

Composing functions

- Function closures can be used to compose two or more functions.

```

(* Use closures to compose two functions *)
(* Type:
   ('b -> 'c) * ('a -> 'b) -> ('a -> 'c) *)
fun compose (f, g) = fn x => f(g x)

(* SML implementation, where the "o" composes f and g *)
f o g

```

Currying

- Functions take “multiple” arguments encoded in one argument via a tuple.

- **Currying** is a closure idiom that takes one argument, and returns a function that takes the next argument, and so on.
 - Currying allows one to call functions using multiple arguments with spaces instead of a tuple expression.
 - SML also has built in syntax support for currying to make implementation less verbose.
 - Function parameters are listed separately with spaces instead of with commas, and aren't enclosed with parentheses.

```
(* Old way to get the effect of multiple arguments *)
fun sorted3_tuple (x, y, z) =
  z >= y andalso y >= x1

val t1 = sorted3_tuple (7, 9, 11)

(* New way using currying *)
val sorted3_curry =
  fn x => fn y => fn z => z >= y andalso y >= x

val t2 = (((sorted3_curry 7) 9) 11)

(* Currying cleaned up with syntactic sugar *)
val sorted3 x y z =
  z >= y andalso y >= x

val t3 = sorted3 7 9 11
```

- It is important to remember that despite the similar syntax to tuples, the currying is semantically returning functions nested within other functions.
 - The syntax support is just syntactic sugar.

Partial application

- If the caller provides too few arguments to a curried function, a function closure is returned that is “waiting for the remaining arguments”.
- This process is called a **partial application**, and can be used to “save” a function for later midway through its evaluation.

```
(* List.filter usually takes two arguments, but by partial application
   one can create a more specific filter function that can be applied
   multiple times *)

val remove_zeros = List.filter (fn x => x <> 0)

(* Produces [1, 3, 7] *)
val list_without_zeroes = remove_zeroes [1, 0, 3, 0, 0, 7]
```

- One may run into the **value restriction** when using partial application on curried polymorphic functions.

- This is a limitation of SML programming language itself.
- This only happens when the resulting function of a partial application is polymorphic.

Currying wrap-up

- Higher-order wrapper functions can be used to:
 - Modify the order of arguments in curried functions.

```
(* Modify the order of arguments in a two argument curried function *)
fun other_curry x y = y x
```

- Convert tupled functions to curried functions.

```
(* Convert tupled function to curried function *)
(* Tupled function *)
range (i, j) =
    ...

(* Curry wrapper function *)
fun curry f x y = f (x, y)

(* Pass "range" into "curry" to create a range function that uses
   currying syntax *)
val range_curry = curry range
```

- Convert curried functions to tupled functions.

```
(* Convert curried function to tupled function *)
(* Curried function *)
range i, j =
    ...

(* Un-curry wrapper function *)
fun uncurry f (x, y) = f x y

(* Pass "range" into "uncurry" to create a range function that uses
   tupled syntax *)
val range_tuple = uncurry range
```

- Currying and tupling multiple arguments are both constant-time operations, so it generally doesn't matter which method is used.
- One should prioritize programming in a way that is elegant, short, concise, and correct.

Mutable references

- Mutable data structures are useful in certain situations.
 - For example, when the state model needs to be updated.

- SML does this with a separate language construct called a `t ref`, where `t` is a type.


```
(* Create a reference with initial contents e *)
ref e

(* Update reference contents *)
e1 := e2

(* Retrieve reference contents *)
!e
```

- A variable bound to a reference is still immutable, as in it will always refer to the same reference.
 - The difference is that the contents of the reference may change via `:=`.
- Aliases to the reference are allowed.

```
val x = ref 42
val y = ref 42
val z = x
val _ = x := 43
val w = (!y) + (!z) (* 85 *)
(* x + 1 does not type-check *)
```



- In the example above, the `val _ =` is necessary because the `x := 43` is an assignment operation. The value name is `_` because that value is not being used anywhere.

Callbacks

- **Callbacks** are when a library takes a function to apply later, when an **event** like a keypress, mouse movement, or data arrival occurs.
 - Different callbacks may need different private data with different types.
 - Fortunately, a function's type does not include the types of bindings in its environment.
- Because of these properties, closures are needed for this language idiom.
- Mutable state is appropriate here, because we want the library to internally maintain a list of registered callbacks that can be updated.

Standard-library documentation

- Many programming languages have a standard library. These libraries exist:
 - For things you cannot implement on your own using the programming language, like opening files or setting timers.
 - For things so common, a standard definition is appropriate, like string concatenation or mapping functions.

- The documentation for Standard ML is:

<https://smlfamily.github.io/Basis/manpages.html>

- REPLs can be used for reminding oneself of how a function works.
 - It gives information about function types, arguments, and sometimes structure.
 - Some programming languages have special support for printing documentation.

▼ Lecture 4:

01/16/24

What is type inference

- **Static type checking** can reject a program before it runs to prevent errors.
 - **Statically typed** programming languages have static type checking.
 - SML is one such statically typed language, where every binding has one type that is determined the program is compiled.
 - SML is **implicitly typed**, meaning that one rarely needs to write down types.
- **Dynamically typed** programming languages do little such checking.
 - Dynamically typed languages determine types during runtime.
- **Type inference** is the giving every binding/expression in a program a type such that type-checking succeeds.

ML type inference

- The key steps for type checking are as follows:
 1. Determine types of bindings in order.
 - a. One cannot use bindings from environments later in the program that the type checker has not gone over yet.
 2. For each `val` or `fun` binding:
 - a. Analyze the definition for all necessary facts.
 - i. For example, `x` in `x > 0` must have type `int`.
 - b. A `TypeError` is raised if there is no way for all facts to hold.
 - i. This would mean that the expression is over-constrained.
 3. Afterward, use polymorphic type variables like `'a` for any unconstrained types.
 4. Finally, enforce the value restriction.

```

val x = 42 (* val x : int *)

fun f (y, z, w) =
  if y (* y must be bool *)
  then z + x (* z must be int *)
  else 0 (* both branches have same type *)
(* f must return an int
   f must take a bool * int * ANYTHING
   so val f : bool * int * 'a -> int
   *)

```

- A central feature of SML type inference is that it infers polymorphic type variables when it can.
- However, type inference and type variables are two separate concepts.
 - Languages can have type inference without type variables, and vice versa.

Type inference examples

- SML statically type checks a function using the rules listed above in order.

```

(*)
1. Determine types of bindings in order:
  f : T1 -> T2 [must be a function; all functions take 1 argument]
  x : T1

  y: T3
  z: T4

2. For each val and fun binding analyze the definition for all facts:
  T1 = T3 * T4 [else pattern match does not type-check]
  T1 = int [abs has type int -> int]
  T4 = int [because we added z to an int]
  so: T1 = int * int
  so: (abs y + z) : int, so let-expression : int, so body : int,
  so: T2 = int

  f : int * int -> int

3. n/a
4. n/a
*)
fun f x =
  let
    val (y, z) = x
  in

```

```
(abs y) + z
end
```

01/17/24

Polymorphic type inference examples

- The type checker infers the existence of polymorphic type variables when a function is under-constrained.
 - All the `T` variables are replaced consistently with type variables like `'a`, `'b`, etc.

```
(*
1. Determine types of bindings in order:
  length : T1 -> T2
  xs : T1

  x : T3
  xs' : T3 list

2. For each binding analyze the definitions for each fact:
  T2 = int
  that's all...

  T3 list -> int

3. Use polymorphic types for unconstrained types:
  T3 = 'a

  'a list -> int
*)
fun length xs =
  case xs of
    [] => 0
  | x:xs' => 1 + (length xs')
```

The value restriction and other type-inference challenges

- SML type-inference is too lenient, and therefore the **value restriction** limits where polymorphic types can occur.
- In the example case below, the type checker would allow `1` to be added to a `string`.
 - This is because the type of `r` is declared to have a polymorphic type.
 - Because `r` is a `ref` and can be mutated, the type can be made more specific in an environment further down, breaking the type checking rules.


```
val r = ref NONE (* val r : 'a option ref *)
val _ = r := SOME "hi"
val i = 1 + valOf (!r)
```

- To restore correctness, there needs to be a stricter type system.
- The **value restriction** solves with by making sure that the only expressions that can have polymorphic types are variables or values.
 - These unconstrained types are filled with dummy types that are essentially unusable.
- Type inference can be done without polymorphism, but this can lead to more restrictive functions.
 - For example, a function that determines the length of a list would only be able to take a list of one specific type.

Mutual recursion

- **Mutual recursion** is when two or more functions can call each other.
 - For example, function `f` can call function `g`, and `g` can call `f`.
- This is useful for multiple common programming idioms and patterns.
- In SML, the keyword `and` can be used to define mutually recursive functions.
 - Every function in the expression are type-checked together and can therefore refer to each other.

```
(* Functions f1, f2, and f3 are allowed to call eachother *)
fun f1 p1 = e1
and f2 p2 = e2
and f3 p3 = e3
```

- This can also be used for defining mutually recursive datatypes.
- Mutually recursive functions can be used for a **state-machine** idiom.
 - A state machine processes an arbitrary length list.
 - As the state machine processes the list, it is always in one of a known number of states.
 - Each next input tells the state machine what “state” to go to next.
 - This can be implemented with mutually recursive function by having one function for each “state”.

Modules for namespace management

- For larger programs, one top-level sequence of bindings is poor style.
- SML provides the keyword `structure` to define **modules** that can separate bindings from other modules.
 - Inside a module, one can use earlier bindings as usual.
 - Outside a module, one must refer to earlier modules’ bindings using `ModuleName.bindingName`.
 - It is convention to capitalize module names.

- SML has built-in modules like `String` and `List`.
- Structure names are not variables, one can only use the bindings inside of them.

```
(* Create a module "MyModule" *)
structure <Name> = struct <bindings> end

(* Refer to module binding *)
val x = <Name>.<binding>
```

- This module system allows for **namespace management**.
 - It gives a hierarchy to names, and therefore avoids shadowing.
 - It allows for different modules to reuse names.

Signatures and hiding things

- **Signatures** are simply types for modules.
 - They describe what bindings a module has and what those bindings types are.
- One can define a signature with `signature` and ascribe it to a module with the `:>` operator.
 - The module will not type-check unless it matches the signature, meaning it has all the bindings in the signature with all the correct types.

```
(* Define a signature *)
signature MY_MATH_LIB
sig
  val fact : int -> int
  val half_pi : int
end

(* Define a structure *)
(* Type:
  structure MyMathLib:
    sig
      val fact : int -> int
      val half_pi : int
    end
*)
structure MyMathLib :> MY_MATH_LIB
struct
  fun fact x = ...
  val half_pi = Math.pi / 2
end
```

- Signatures are useful for **hiding bindings** and type definitions.
- It is useful to have a “private” top-level function so that two or more functions could easily share a helper function.

- SML does this with signatures that **omit bindings**.
- One defines private bindings that can only be used in the module by omitting that binding from the module's assigned signature.

```
(* Omit half_pi from the signature, essentially making it a private binding *)

(* Define a signature *)
signature MY_MATH_LIB
sig
    val fact : int -> int
end

(* Define a structure *)
(* Type:
    structure MyMathLib:
        sig
            val fact : int -> int
            val half_pi : int
        end
*)
structure MyMathLib :> MY_MATH_LIB
struct
    fun fact x = ...
    val half_pi = Math.pi / 2
end

(* This will not work, as half_pi cannot be accessed outside of MyMathLib *)
val x = MyMathLib.half_pi
```

Abstract data types using signatures

- Modules should hide concrete type definitions so clients cannot create invariant-violating values of the type directly.
 - This is achieved with **abstract data types (ADTs)**, which are custom types that exist but have inaccessible definitions.
 - One can define ADTs in signatures with the keyword **type**.

```
(* Define a type "natural" which can be type checked *)
(* This type "natural" cannot be accessed outside the structure in which
    the signature MY_MATH_LIB is used *)

signature MY_MATH_LIB
sig
    type natural
    (* fact takes an ADT natural, and returns int *)
    val fact : natural -> int
end
```

- As a result, there are two ways to use signatures to hide data from an external user of a module.
 1. Omit bindings from the signature so they cannot be used outside a module.
 2. Create abstract data types so that values cannot be created from them outside the module.

Signature matching

- **Signature matching** is what the type checker does to check if a `structure` can have a `signature`.
- The type checker checks if every non-abstract datatype, abstract datatype, `val` binding, and exception in the signature is provided in the structure/module.
 - The module can be more general than the signature.
 - The signature must be more specific than the module.

Equivalent implementations

- **A key purpose of abstraction is to allow different implementations to be equivalent.**
 - In essence, a function should be a “black-box”, where the end user does not need to know how it works.
 - It allows the implementation of a function to be improved or changed as long as the functionality doesn’t change.
 - This is easier to do if one starts with more abstract signatures that reveal less to an end user.

Different modules define different types

- Although multiple signatures can define abstract datatypes with the same name, each implementation introduces a new type.
- This means that when passing a module’s unique type into another module containing a type of the same name, it does not type check.

```
(* Although the types may seem to be equivalent, natural in MathLib1
   actually has type MathLib1.natural, and natural in MathLib2 actually
   has type MathLib2.natural
```

```
signature MLIB_1
sig
  type natural
end
structure MathLib1 :> MLIB_1
struct
  ...
end
```

```
signature MLIB_2
sig
  type natural
end
structure MathLib2 :> MLIB_2
```

```
struct
  ...
end
```

Equivalent functions

- Understanding when two pieces of code are equivalent is a fundamental software engineering idea.
 - Equivalence is made easier with more abstraction and fewer side effects.
- Code maintenance, backward compatibility, optimization, and abstraction rely on Equivalence.
- This raises the question of how one can define that two functions are equivalent.
- Two functions are equivalent if they have the same “observable behavior” no matter how they are used anywhere in a program.
 - They have the same input and the same output.
 - They produce equivalent results.
 - They raise the same exceptions.
 - They have the same termination and non-termination behavior.
 - They mutate non-local memory in the same way.
- Pure functions are functions without side effects.
 - Pure functions are therefore preferred more maintaining function equivalence.

Standard equivalences

- Syntactic sugar is, by definition, always equivalent with the semantics that it is syntactic sugar for.
 - If an implementation changes the evaluation order of a program, it cannot be syntactic sugar.
- Consistently renamed bound variables maintains function equivalence.
- Using a helper function maintains function equivalence.

Equivalence versus performance

- Even though functions are equivalent, that doesn’t mean that they are just as efficient.
- Program language equivalence refers to having the same inputs, outputs, and side effects.
- Asymptotic equivalence refers to focusing on time efficiency for an algorithm for a large input.
- Systems equivalence is like asymptotic equivalence, but accounts for differences between time efficiencies with different coefficients.
 - This equivalence does not study the general algorithm like asymptotic equivalence does.
- All three of these equivalences are important tools that are used to help software engineers write efficient programs.