

Programming Languages, Part C

▼ Status

Completed

▼ Lecture 8:

01/24/24

Introduction to Ruby

- Ruby is a **pure object oriented language (OOP)**, meaning all values are objects.
- Ruby is **class-based**, meaning every object has a class that determines its behavior.
- Ruby is a *very* dynamically typed programming language.
- Ruby is a modern “scripting language” popularly used for building server-side web applications.

```
# Use "ruby" to run a ruby program
$ ruby "program.rb"

# Type "irb" (Interactive Ruby) to run the Ruby REPL
$ irb

# Use "load" to run a ruby program in the REPL
$ load "program.rb"

# Type "quit" to exit the REPL
$ quit
```

Classes and objects

- There are rules that guide one’s understanding of the Ruby programming language:
 1. All values are reference to objects.
 2. Objects communicate through **method/message** calls.
 - a. Methods are functions that belong to objects.
 3. Each object has its own private state.
 4. Every object is an instance of a class.
 - a. An object’s class determines the object’s behavior.
 - i. How the object handles method calls.
 - ii. The class contains method definitions.
- New classes are defined with methods, and objects can be initialized from those class definitions.
 - New objects can be created with the syntax `ClassName.new`.
 - Object methods can be called with the syntax `object.method_name(args)`.

```
# Anatomy of a Ruby class
```

```
class <ClassName>
  def <method_name> <method_args>...
    <expression>
  end
  ...
end
```

```
# Example "Hello" class that contains a method that prints "Hello, World"
```

```
class Hello
  def hello_world
    # "puts" is a method that prints a string
    puts "Hello, World!"
  end
end
```

```
# Create an object "x" whose class is the class "Hello"
x = Hello.new
```

```
# Evaluate "x" to an object, and call its "hello_world" method
x.hello_world
```

- `self` refers to the object *itself*.
 - Or more specifically, the current object that the method is called on.
 - One can pass, return, or store the whole object with the syntax `self` by itself.

```
# The method m2 will call its object's method m1, resulting in x + y + 34
# It sends the same object an "m1" method
```

```
class A
  def m1
    34
  end
  # method arguments are surrounded by parentheses and separated by commas
  def m2 (x, y)
    # Syntactic sugar for "self.m1" is just "m1"
    x + y + self.m1
  end
end
```

- Methods can use local variables, and their scope is the method body.

- There is no need to declare methods or variables, as they can be assigned anywhere.
 - Variables are mutable, and are also allowed a the “top-level” global scope.
 - The contents of a variable are always a reference to an object.
- In Ruby, new lines affect semantics, but indentation does not.

Object state

- In Ruby, all objects have **state**.
 - The object’s state persists, as in it can grow and change from the time the object is created.
 - An object’s state is only directly accessible from an object’s methods.
 - An object’s state consist of **instance variables**, commonly known as **fields**.
 - In Ruby, fields are assigned with `@`.
 - If an object tries to access an undefined field, it is evaluated to a `nil` object.
- Creating a new object always returns a reference to a new object.
- Because mutation exists, one should be careful with **aliasing**.
 - Variable assignment creates an alias, meaning multiple variables hold references to the same object.
 - Any changes to an object through one variable affect that same object, for all the other alias variables to that object.

```
# Create two distinct Counter objects
a = Counter.new
b = Counter.new

# Create two aliases to the same Counter object
a = Counter.new
b = a
```

- The built-in method `initialize` is called when objects are created.
 - `initialize` will be called before `ClassName.new` returns the object being created.
 - Arguments can be passed with `new` to the `initialize` method.
- This is not, like it is in many programming languages, a constructor.
 - In Ruby, one is not required to initialize fields in `initialize`, and it is only convention.
- **Class variables** are state that is shared by the entire class.
 - These variables are shared, and only accessible, by instances of the class.
 - These are syntactically written with `@@`.
- **Class constants** are variables that syntactically start with a capital letter.
 - These class constants are allowed to, but should not, be mutated.
 - They are accessible outside the class with `ClassName::Constant`.
- **Class methods**, also known as **static methods**, are written with the prefix `self.`.

- Class methods are called with `ClassName.method_name(args)`.
- These methods are part of the class, not a particular object instance of it.

Visibility

- Hiding expressions and values is essential for modularity and abstraction.
- In Ruby, object state is always **private**, as in instance variables can only be accessed through an object's methods.
 - To make object state publicly visible, one should define **getters** and **setters**.

```
# Get the variable foo from the object
def get_foo
  @foo
end

# Set the variable foo for the object
def set_foo x
  @foo = x
end
```

- Ruby support syntactic sugar for these getters and setters.

```
# Get the variable foo from the object
def foo
  @foo
end

# Set the variable foo for the object
def foo= x
  @foo = x
end

# Call foo= two different ways
e.foo= y
# " =" is a method call to "="
e.foo = y
```

- Requiring private instance variables forces indirect interfacing with an object, without knowing about the actual implementation of the object.
 - This is essential for abstraction and modularity, as it allows class implementation to change without introducing bugs.
- There are three visibilities for Ruby methods:
 1. **private** methods are only available to the object itself.
 - a. If a method is private, one can only call it with the shorthand `method_name(args)` and never with `self.method_name(args)`.
 - i. This shorthand does not indicate what object the method is being called on.

2. `protected` methods are only available to objects of the same class or subclasses.
3. `public` methods are available to the whole program.
 - Methods are `public` by default.

01/25/24

Everything is an object

- Ruby is fully committed to object oriented program, meaning every value is a reference to an object.
 - This leads to simpler, smaller semantics.
 - The **only operation** on objects is calling methods on them.
 - All top-level methods are just added to the built-in `Object` class.
 - This `Object` class contains all built-in methods in Ruby.
 - All classes defined are **subclasses** of `Object`.
 - Therefore, all `Object` methods are passed down and able to be used by any object.
- Every class has an “undefined” method that is called when a method that doesn’t exist is called for an object.
- `nil` is an object that is similar to the `null` value in Racket.
- All objects have built-in methods like `methods` and `class`.
 - `methods` returns an object’s methods.
 - `class` returns an object’s class
 - This process of querying information about an object during runtime is called **reflection**.
 - This can be useful in the REPL to explore what methods are available without having to consult the Ruby documentation.

Class definitions are dynamic

- Ruby programs can add, change, or replace methods while a program is running.
 - This can break abstractions and make programs difficult to analyze.
 - This helps re-enforce the rules of OOP.
 - Every object has a class, and a class determines its instances’ behavior.
 - Since a class in an object itself, changes to a class will reflect those changes on its object instances.

```
# Define a class Rational
class Rational
  def initialize x
    @value = x
  end
  ...
end

# Create an object with the class Rational
```

```

a = Rational.new

# Modify the class Rational
class Rational
  def double
    self + self
  end
end

# The object has not changed, but its class has
# Therefore, the object can now call the method "double"
a.double

```

- This behavior can even be used to modify Ruby's built-in classes.
 - Any new method defined at top-level is added to the `Object` class.
- Dynamic features like these can create interesting semantic questions that must be answered.
 - More dynamic features can lead to more questions that need to be answered, decreasing a programming language's performance.

Duck typing

“If it walks like a duck and quacks like a duck, it's a duck”

- When one needs to pass an object into a method, in reality one usually only needs to pass in an object with the required certain properties (methods) of that object.
 - **Duck typing** means passing an object with the needed methods, instead of the entire original object itself.
 - It “walks” and “quacks” like the original method, and therefore for all practical purposes, it is the original method.
 - One advantage is that this allows for more code reuse, since methods an object receives is “all that matters”.
 - One disadvantage is that abstraction is not as useful, since one cannot change the implementation of methods that assume duck typing.

```

# Example of "duck typing"

# The naive perspective:
# Method takes in a point and mirrors it along the x-axis

# The duck typing perspective:
# In reality it works for
# anything with x= and x (whose result has a * method), methods

def mirror_update point
  point.x = point.x * -1
end

```

01/26/24

Arrays

- The most commonly used data structure in Ruby is the `Array` class.
 - An array holds any number of other objects, and is `indexed` by number, starting at 0.
 - Similar to Python, Ruby interprets negative numbers as counting from the end of an array.
 - Array size is not fixed, and elements can be dynamically created or removed during runtime.

```
# Initialize an array object
a = [1, 2, 7, "Hello, World!", -1]
# Initialize an array of length 7, where every element is nil
b = Array.new(7)
# Initialize an array of length 10, where every element is "hi"
c = Array.new(10) { "hi" }

# Bind the "i"th index of "a" to "d"
d = a[i]

# Set the "i"th index of "a" to "e"
a[i] = e
```

- These arrays are much more flexible and dynamic (but less efficient) than in other programming languages.
- Arrays can be used as `stacks`, since they have the operations `push` and `pop`.
 - `push` adds an element to the end of an array.
 - `pop` removes and gets the last element of an array.
 - `unshift` adds an element to the beginning of an array.
 - `shift` removes and gets the first element of an array.
- Arrays, just like any other objects, can be aliased.
- Arrays can be `sliced`.

```
# Create a new array
a = [1, 2, 3, 4, 5, 6]

# Assign a slice of the array to b,
# starting at element 2 and giving 3 elements
# [3, 4, 5]
b = a[2, 3]
```

Blocks

- `Blocks` are a feature in Ruby that allow any easy way to pass anonymous functions to methods.
 - Blocks can take zero or more arguments.
 - Blocks use lexical scope, where the block body uses the environment where it was defined.

- Syntactically, one can use `{ ... }` or `do ... end` to enclose a block.

```
# Anatomy of a block
{ | args... | expression }
# Alternatively replace { } with do end
do | args.. | expression end

# Run a block 3 times
# Prints "Hello, World!" 3 times
3.times { puts "Hello, World!" }

# Call the block once for each element of an array
# Print the square of each element in the array
[1, 2, 3, 4].each { |x| puts (x * x) }

# Create an array of the first 5 multiples of 4, starting at 4
Array.new(5) { |i| 4 * (i + 1) }
```

- One can pass zero or one blocks with *any* message (method).
- Blocks can be used to apply higher-order functions to arrays.
 - Because blocks are so useful for applying higher-order functions to arrays, loops are very rarely used.

Using blocks

- When a callee uses a block argument, there is no name for that block argument.
- When defining the use of blocks inside a method, the keyword `yield(args)` is used to denote where a block is used, and what arguments the block uses.
 - If a caller tries run a method that contains `yield` without a block, an error will be raised.
 - `block_given?` evaluates to `true` if a block is passed by a caller, and `false` otherwise.

```
# Define the method silly which takes a block
def math_op
  yield(4, 5) + yield(100, 100)
end

# Call the method silly with a block
# Yield evaluates its arguments using the block it is passed
# Result is 103 = [(2 * 4) - 5] + [(100 * 2) - 100]
math_op { |a, b| (2 * a) - b }
```

Procs

- **Procs** are similar to blocks, except they are actual objects with the abilities of function closures.
 - Blocks are “second-class” expressions, as their only function is to be yielded too.
 - They cannot be returned, passed, stored, etc.

- Instances of the `Proc` class are “first-class” expressions.
 - **First class** indicates that an expressions can be the result of a computation, returned, stored, and so on.
- Using the method `lambda` on an `Object` object takes a block and returns its corresponding `Proc`.

```
# Trying to create an array of closures is not possible with blocks
# A block cannot pass in another block

# This would return an error
c = a.map { |i| { |y| i > y } }

# Trying to create an array of closures is possible with Procs

# The lambda method of Object is used to turn
# the block into an instance of Proc

# Because lambda is a method of the Object class, and all top-level functions
# are part of the Object class, there is no need to write Object.lambda
c = a.map { |i| (lambda { |y| i > y }) }
```

- First-class expressions make closures more powerful than blocks.
 - In simple cases, blocks can be more convenient to use.

Hashes and ranges

- **Hashes** and **ranges** are two standard classes in Ruby that are commonly used in Ruby programs.
 - Hashes are arrays, with mappings from keys to values.
 - The keys can be any type of object, and as a result there is no natural ordering of numeric indices that would be found in an array.

```
# Create an empty hash
hash = {}

# Assign the value "Nate" to the key "First Name"
# {"First Name"=>"Nate"}
hash["First Name"] = "Nate"

# Assign the value "Levine" to the key "Last Name"
# {"First Name"=>"Nate", "Last Name"=>"Levine"}
hash["Last Name"] = "Levine"

# Return an array of the hash's keys
# ["First Name", "Last Name"]
hash.keys

# Return an array of the hash's values
```

```
# ["Nate", "Levine"]
hash.values
```

- Ranges are arrays of contiguous numbers that are more efficiently represented than arrays.

```
# Create and range object from 1 to 100
range = (1..100)

# Add up all the numbers from 1 to 100 (5050) using the method inject
range.inject { |acc, elt| acc + elt }
```

- It is generally good style to:
 - Use ranges instead of arrays instead of numbers where one can.
 - Use hashes instead of arrays when non-numeric keys are better suited for representing data.

Subclassing

- Sub-classing is a concept essential to object oriented programming.
- A class definition always has a **super-class/parent class**.
 - This super-class is **Object** by default.
 - A class **inherits** all the method definitions of its super-class.
 - A class can **override** method definitions as desired.
- Unlike Java/C#/C++, classes in Ruby cannot inherit fields, since all instance variables in Ruby are not part of class definitions.
- The **super** method is used inside a sub-class's method to call a method of the same name from the super-class, that has been overridden by the sub-class.

```
# The superclass Point
class Point
  def initialize(a, b)
    @x = a
    @y = b
  end

  def distFromOrigin
    Math.sqrt(@x * @x + @y * @y)
  end
end

# The subclass ColorPoint
# The < denotes it is a subclass of the class Point
# ColorPoint inherits the method "distFromOrigin" from Point

# ColorPoint's initialize method overrides Point's initialize method
```

```
# "super" is used to call the superclass' method of
# the same name as a helper function
class ColorPoint < Point
  def initialize(x, y, c="clear")
    super(x, y)
    @color = calls
  end
end

# Get the superclass of ColorPoint (Point)
ColorPoint.superclass
```

- Sub-classes *are the same classes* as their super-classes.
 - Instances of sub-classes *are not instances* of their super-classes.

```
# ColorPoint is a Point
# true
Point.is_a? ColorPoint

# ColorPoint and Point are Object
# true
ColorPoint.is_a? Object
# true
Point.is_a? Object

# ColorPoint is not an instance of Point
# false
ColorPoint.instance_of? Point
```

- Using methods like `is_a?` and `instance_of?` is usually not OOP style, as it disallows techniques like duck typing where the class of an object can be arbitrary.

Why use subclassing?

- Sub-classing is a useful programming feature, though it tends to be overused in OOP programming languages.
- Sub-classing removes the need to add new methods to a super-class.
 - Adding new methods to super-class can break modularity.
- Sub-classing removes the need to have the same methods in multiple similar classes.
 - Creating subclasses that inherit from super-classes reduces code reuse.
- An alternative to sub-classing is having an instance of a super-class that initialized inside a “sub-class”.
 - Unfortunately, this is less convenient code reuse
 - In addition, the sub-class would not be the same class as the super-class.

- This means that the sub-class cannot be used where a program expects the super-class.

```
# Instead of making ColorPoint a sub-class of Point,
# Initialize a Point object inside the ColorPoint and forward
# messages to it through ColorPoint
class ColorPoint
  def initialize(x, y, c="clear")
    @pt = Point.new(x, y)
    @color = c
  end
  def x
    @pt.x
  end
  def y
    @pt.y
  end
end
```

Overriding and dynamic dispatch

- With the examples so far, one can argue that objects are not so different from closures.
 - Unlike closures, objects have multiple methods.
 - Unlike closures, objects have instance variables rather than an environment.
 - Object inheritance avoids helper functions or code copying.
- The big difference is that overriding can make a method defined in the super-class call a different method in the sub-class.
 - Overriding a method in a sub-class is sometimes necessary to maintain correct logic for that method.
- When one has a method that makes another call on the same object (when that method uses `self`), that `self` is the entire object.
 - If `self` is an instance of the sub-class, then the sub-class's methods are used, *not* the super-class's methods.

Method-lookup rules

- To understand dynamic dispatch fully, it is important to precisely understand the semantics of method look-up in OOP.
- **Dynamic dispatch** (also known as **late binding** or **virtual methods**) refers to the semantics that when calling `self` on a method `m1` in another method `m2` in class `c`, `m2` can **resolve** to a method `m2` defined in a sub-class of `c`.
 - This is the most unique characteristic of OOP that distinguishes it from closures or functions.
- All lookup rules in Ruby are defined in terms of `self`.
 - `self` maps to some “current” object.
 - Instance variables `@x` are looked up in the object bound to `self`.
 - Class variables `@@x` are looked up in the object bound to `self.class`.
 - For methods calls:

1. All expressions called with a method are evaluated to objects.
 2. Let `c` be the class of the expression that the method will be called on.
 - a. `c` is the **receiver** of the method.
 3. If the method is defined in `c`, use that method.
 - a. Otherwise, check the super-class of `c` for that method, and that super-class's super-class, and so on until the method is found or `Object` is reached.
 - i. If the `Object` class is reached, a `method_missing` error message is raised by it.
 - b. Whatever class the method is found in first, that method from that class is called.
 4. The body of the method picked is evaluated.
 - a. `self` is bound to the expression that receives the method.
 - i. This is the key point that implements dynamic dispatch.
- A method body is evaluated in an environment where `self` refers to the object that the method is called on.

```
# obj.b evaluates to "Hello, World!"
# 1. When obj calls method "b", method "a" is called
# 2. The current object "obj" has class B
# - Because class B contains no method "a",
#   "obj" searches its super-class A
# 3. method "c" is called in class A, and that method calls method "b"
# 4. Because the current object (the object that received method "a")
#   has class B, the method "c" in class B is called,
#   not the method "c" in class A
# 5. Method "c" in class B evaluates to "Hello, World!"
obj = B.new
obj.b

# Class A is a super-class of class B
class A
  def a
    self.c
  end
  def c
    42
  end
end

# Class B is a sub-class of class A
class B < A
  def b
    # Method "a" will evaluate using the object with class B as "self"
    self.a
  end
  def c
```

```

    "Hello, World!"
  end
end

```

- Although these semantics for methods are more complex than function closures, it does not mean they are inferior or superior to them.

▼ Lecture 9:

01/31/24

OOP versus functional decomposition

- When comparing OOP and functional programming, one should consider how these two programming methods decompose problems into more manageable pieces.
 - Functional programming breaks problems down into functions that each perform their own operation.
 - OOP breaks programs down into classes that give behaviors to data.

	<code>eval</code>	<code>toString</code>	<code>hasZero</code>	...
<code>Int</code>				
<code>Add</code>				
<code>Negate</code>				
...				

- Functional programming essentially breaks down to applying operations to variants (types of data)
 - This application of operations to variants can be summarized in a table.
 - Each column is evaluated through a case expression, therefore **the code is organized primarily by the columns (operators)** of the “grid”.

```
(* Example of function-based organization in ML *)
```

```
(* The functions are organized by operations, and each function
   uses a case expression to evaluate different variants *)
```

```
fun toString e =
  case e of
    Int i => ...
  | Add (e1, e2) => ...
  | Negate e1 => ..
```

- OOP defines classes with one **abstract method** for each operation, and subclasses are defined for each variant.
 - Each row in the “grid” represents a class with one method implementation for each “grid” column.
 - This is essentially the opposite of functional programming, as in this case **the code is organized primarily by the rows (classes)** of the “grid”.

```
# Example of class-based organization in Ruby

# The classes are organized by variants, and each class has methods
# that operate on expressions
class Int < Exp
  def Int
    ...
  end

  def Add
    ...
  end

  def Negate
    ...
  end
end
```

- In a sense, **functional programming and object oriented programming are doing the same thing in the exact opposite way.**
 - The question “boils down to” whether one wants to organize their program by rows or by columns.

Adding operations or variants

- Planning for software extensibility can influence which programming style is used.
 - If a new variant will need to be added to a program, OOP style may be used to add another row to the “grid”.
 - If a new operation will need to be added to a program, functional style may be used to add another column to the “grid”.
- Making extensible software is often difficult yet valuable, as code can be reused more.
 - Some modern languages, such as Scala, attempt to support extensibility in both directions.

Binary methods with functional decomposition

- Often operations are defined over multiple variants.
- For example, one may want to redefine an **Add** operation to work over the variants **Int**, **String**, and **Rational**.
 - Addition is a **binary** method/operation, because it takes two variants to evaluate.
 - This operation needs to be redefined using another “grid”.
 - In functional implementation, this is achieved through a helper function.

```

(* Example of function-based organization in ML with a binary operation *)

(* Helper function which decides the evaluation of a binary method "Add" *)
fun add_values (v1, v2) =
  case (v1, v2) of
    (Int i, Int j) => ...
  | (Int i, String s) => ...
  | (Int i, Rational (i, j)) => ...
    ...
  | (String s1, String s2) => ...
  | (String s, Rational (i, j)) => ...
  (* Else case covers the case where v1 or v2 are
     not Int, String, or Rational *)
  | _ => ...

fun toString e =
  case e of
    Int i => ...
  (* A helper function is used to decide the evaluation of Add *)
  | Add (e1, e2) => add_values (eval e1, eval e2)
  | Negate e1 => ..
  | String s => ...
  | Rational i/j => ...

```

	Int	String	Rational
Int			
String			
Rational			

Double dispatch

- In OOP, there are no case functions to take care of binary operations.
 - This means that one variant needs to call a binary method on another variant.
- When this happens, it is impossible to know the class of the variant `v` being passed into the binary method.
 - To solve this problem, a method must be called on `v` to tell it what variant of the receiver of the method `self` is, by passing `self` into that method.
 - This programming idiom is called **double dispatch**.


```

# The class Int has a method "add_values" which takes a variant v
# and adds itself (an Int) to v

# It does this by passing itself into a method v (which could be an Int,
# String, Rational, etc.), whose class has a method "addInt" which
# decides how it adds and Int to itself

class Int
  # First dispatch
  def add_values v
    v.addInt self
  end

  # Dispatches for when Int is the variant passed into an "add_values" method
  def addString v
    ...
  end
  def addNatural v
    ...
  end
  ...
end

class <Type of v>
  # Second dispatch
  def addInt v
    ...
  end

  # Dispatches for other variants
  def addString v
    ...
  end
  def addNatural v
    ...
  end
  ...
end

```

- Every class has the appropriate “cases” (not actually cases like in SML, but methods) to handle how they evaluate binary methods with both their class and another class.

Multi-methods

- An alternative to double dispatch is [multiple dispatch/multi-methods](#).
 - For example, each class would have multiple methods of the same name, each one corresponding to an operation on a different class.

- Each method indicates what instance of what class it takes as an argument.
- However, interactions of multiple dispatch with subclassing can confuse an end user over which methods should be called for what situations.
- This is not an appropriate idiom for Ruby, as Ruby:
 1. Places no restrictions of what arguments can be passed into a method.
 2. Does not allow multiple methods with the same name.

Multiple inheritance

- It is important to distinguish the relation between different classes with more specific terminology than just “sub-class” or “super-class”.
- Class hierarchies form a tree:
 - Each node is a class.
 - A sub/super-class is an **immediate** sub/super-class when they are **separated by one level** of inheritance.
 - **Parents** are immediate super-classes.
 - A sub/super-class is a **transitive** sub/super-class when they are **separated by more than one level** of inheritance.
- **Multiple inheritance** allows for a sub-class to have more than one super-class.
 - Languages like C++ allow this, where other languages like Ruby and Java do not.
 - With multiple inheritance, the class hierarchy is no longer a tree.

```
graph TD
  X --> V
  X --> W
  V --> Y
  W --> Y
```

- In the tree above, a **diamond** is formed when multiple paths show that X is a transitive superclass of Y .
 - This can lead to ambiguity in a program.
 - For example, if both V and W define a method `m`, which method `m` does Y inherit?

Mixins

- Ruby has a feature called **mixins**, which are an alternative to multiple inheritance.
 - Mixins are collections of methods.
 - Mixins are not classes, and therefore cannot be instantiated.
- Ruby allows for a sub-class to have only one superclass, but multiple mixins.
 - Semantically, mixins make their methods part of the class.
 - Mixins are more powerful than helper methods, as mixin methods can access methods on `self`.

- Mixins allow classes to gain many methods for a small amount of work, without the complexity of multiple inheritance.
- Ruby mixins are defined using the `module` keyword.
- Mixins can be included in classes using the `include` keyword.

```
# Example of how the mixin Doubler can be included in Pt
# to give Pt a "double" method

# Mixin "Doubler"
module Doubler
  def double
    # We assume that any class that includes "Doubler" will have a "+" method
    # When "double" is called in Pt, Pt will call the "+" method.
    self + self
  end
end

class Pt
  attr_accessor :x, :y
  # Now if Pt.double is called, a new point with double the original
  # point's x and y coordinates will be returned
  include Doubler
  def + other
    ans = Pt.new
    ans.x = self.x + other.x
    ans.y = self.y + other.y
  end
end

# Return a point with x = 6, y = 8
Pt.double(Pt.new(3, 4))
```

- There are two popular mixins in Ruby, those being:
 1. `Comparable`, which defines `<`, `>`, `==`, `!=`, `>=`, and `<=` in terms of `<=>`.
 - a. `<=>` is sometimes called the “spaceship operator”.
 2. `Enumerable`, which defines many iterators such as `map`, `find`, and so on.

Interfaces

- Another alternative to multiple inheritance and mixins is with Java/C# style `interfaces`.
 - Interfaces rely on static typing, and are therefore not semantically possible in Ruby.
- Static typing works slightly differently for OOP than it does for functional programming.
 - Sound typing for OOP prevents “method missing” errors.

- The type checker will check if objects possess the methods called on them.
- Each class has a type, and each method has argument and result types.
 - Any subclass is also a **subtype**.
 - A subtype can be used anywhere a supertype is allowed.
- **Interfaces** are types, not classes.
 - Interfaces do not contain method definitions, only method signatures.
 - Interfaces are not instantiable.
- A class can implement any number of interfaces.
 - For a class to type check, each method in the interface must have the correct type.
 - If a class type-checks, it is a subtype of the interface.
 - If class A implements interface I , A is a subtype (not a subclass) of I .
- Interfaces are primarily used for making a type system more flexible, so a callee can call certain methods regardless of their class.
- Because interfaces allow what is essentially the default in many dynamically typed programming languages, they have very little use in those languages.

Abstract methods

- One should consider how required overriding is implemented in a statically typed programming language.
 - **Required overriding** refers to when a super-class requires its sub-classes to override one or more of its methods.
- Often a class is written in such a way that it expects its sub-classes to override some methods.
 - The purpose of the super-class in this case is to abstract common functionality.
 - Because the super-class is “abstract”, it does not make sense to make an instance of the super-class.
- Java/C#/C++ let super-classes give a type signature for methods that sub-classes need to provide.
 - These are called **abstract methods** in Java, and **pure virtual methods** in C#/C++.
 - This also disallows instances of the super-class from being created.
- Abstract methods do not make these programming languages more powerful.
 - In fact, their point is to limit the functionality of classes, not expand it.

▼ Lecture 10:

02/02/24

Subtyping from the beginning

- This lecture will only contain pseudocode, as none of the three programming languages studied in this course have subtyping as a feature.
- Most core subtyping ideas can be covered by considering records with mutable fields.
 - Therefore the pseudocode will be most similar to SML, as that programming language has records with *immutable* fields.

- Pseudocode syntax is as follows:
 - Each record field `f` contains an expression `e` that evaluates to `v` and has a type `t`.

```
# Record creation (field = expression)
# evaluate e's, make a record
{f1=e1, f2=e2, ..., fn=en}

# Record field accessing
# Evaluate e to a record v with a field f
e.f

# Record field updating
# Evaluate e1 to a record v1, e2 to a record v2, change v1's field to v2
# e1 and e2 must have the same type t
e1.f = e2

# Each field f has a type t
# e.f has type t
{f1:t1, f2:t2, ..., fn:tn}

# Example value
val point : {x:real, y:real} = {x=3.0, y=4.0}
# Access x and y from point
point.x
point.y
```

- If an expression has type `{f1:t1, f2:t2, ..., fn:tn}`, then **subtyping** would allow it to have some fields removed.

```
# This call to "dist" would only work with subtyping,
# as without it color_pt contains too many fields

fun dist (p:{x:real, y:real}) =
  ...

val color_pt : {x:real, y:real, color:string} = ...
val _ = dist(color_pt)
```

The subtype relation

- Subtyping rules are defined exactly, in a way that doesn't change any typing rules that have already been encountered, while also providing more flexibility to a type system.
 - Programming languages already have a lot of typing rules that would be too complicated to change.
- This can be done by adding two rules to a language:
 1. The subtyping rule, `t1 <: t2`, stating that `t1` is a subtype of `t2`.

2. One new typing rule that uses subtyping.
 - a. If `e` has type `t1` and `t1 <: t2`, then `e` also has type `t2`.
- These rules are not a matter of opinion, since any variation from these subtyping rules may result in an unsound programming language.
 - If a programming language is sound before adding subtyping rules, it should be kept that way.
 - For example, one may be able to access record fields that do not exist.
 - These rules achieve the principle of **substitutability**, meaning that if `t1 <: t2`, then any value `t1` is usable in any way `t2` is.
- These rules follow from the substitutability principle:
 1. **Width subtyping**: A wider record can be a subtype of a slimmer record, as long as the slimmer record has the same types as the wider record.
 2. **Permutation subtyping**: A supertype can have the same fields as a subtype with the same types, just in a different order.
 3. **Transitivity**: If `t1 <: t2` and `t2 <: t3`, then `t1 <: t3`.
 4. **Reflexivity**: Every type is a subtype of itself, so `t <: t`.

Depth subtyping

- **Depth subtyping** suggests that if `ta <: tb`, then a record with a field `tb` can be subtyped with a new record where that field now has type `ta`.

```
# What a depth subtyping rule may look like
if ta <: tb
then {f1:t1, ..., f:ta, ..., fn:tn} <: {f1:t1, ..., f:tb, ..., fn:tn}
```

- Unfortunately, depth subtyping is unsound, as it allows programs to access missing record fields.
 - For example, a function could mutate a record to take a subtype of that record with fewer fields.
 - Now when trying to access one of the fields that has been removed, the program breaks.

```
# sphere is defined with a z field
val sphere:{center:{x:real, y:real, z:real}, r:real} =
  {center={x=3.0, y=4.0, z=5.0}, r=1.0}

# get2dCenter removes the z field from the sphere's
# center's z field through depth subtyping
fun get2dCenter (c:{center:{x:real, y:real}, r:real}) =
  c.center = {x=x, y=y}

val sphere2dCenter = get2dCenter(sphere)

# sphere's center's z field cannot be accessed, as it has no z field anymore
val sphere_z = sphere.center.z
```

- This a good example of how when one creates a rule for a programming language, it cannot just be evaluated it on the programs that one wants to work.
 - One also has to make sure that the programs one *does not want* are still not allowed.
- For depth subtyping to work, record fields would need to be immutable.

02/03/24

Function subtyping

- Subtyping already allows for:
 - Function callers to pass in a subtype of a parameter as an argument.
 - Functions to return the subtype of their resulting value(s).
- One should consider how a programming language defines when a function itself is a subtype of another function.
- Following from the two base subtyping rules, one can derive the rule that if $ta <: tb$, then $t \rightarrow ta <: t \rightarrow tb$.
 - This means that the functions are **covariant** in their return type.
- This *does not* mean that if $ta <: tb$, then $ta \rightarrow t <: tb \rightarrow t$.
 - This rule breaks soundness.
- However, the opposite rule that if $tb <: ta$, then $ta \rightarrow t <: tb \rightarrow t$ is sound.
 - This means that the function arguments are **contravariant**.
- Therefore if $t3 <: t1$ and $t2 <: t4$, then $t1 \rightarrow t2 <: t3 \rightarrow t4$.
 - Function subtyping is contravariant in its arguments, and covariant in its results.

Subtyping for OOP

- Our understanding of subtyping for records are functions can be used to understand subtyping for class-based OOP, and how a static type checker works for OOP.
 - Class names are types, and therefore sub-classes are sub-types.
- An instance of a sub-class should be usable anywhere in place of an instance of its super-class.
- In a sense, objects are essentially records that hold *mutable fields*, and *immutable methods* that have access to `self`.
- It is important to understand the distinction between classes and types.
 - A class defines an object's behavior.
 - Subclassing inherits behaviors and changes that behavior.
 - A type describes an object's methods' argument and return types.
- `self` is special in the fact that it is an exception to the contravariant argument rule for methods/functions.
 - `self` can be used as a covariant argument for class methods.

Bounded polymorphism

- **Generics** are polymorphic type variables, like `'a` in SML.
- While generics and subtyping are applicable for separate use cases, programming languages can have both features.

- These two features can be combined to create **bounded polymorphism**, where any type `'b` can be a subtype of `'a`.
- Subtyping alone cannot be used for certain functions.

```
# This would not be allowed, since List<ColorPoint>
# is not a subtype of List<Point>
# This is due to the fact that depth subtyping is not allowed,
# as in the elements of one list being subtypes of the elements of
# the other list is not considered

List<Point> inCircle (List<Point> pts) =
    ...

# Passing in a list of objects with the class Point is allowed
inCircle(pts)

# Passing in a list of objects with the class ColorPoint is not allowed
inCircle(color_pts)
```

- In addition, generics alone cannot be used for certain functions.

```
# This would not be allowed, since List<T> can pass in
# anything, not just objects with x and y fields

List<T> inCircle (<T> pts) =
    ...
```

- Therefore, it is useful to create a generic that is a subtype of a specific supertype.

```
# This is allowed as long as T is a subtype of Point, or is type Point

# T <: Point
List<T> inCircle (List<T> pts) =
    ...
```

Wrap-up

- This course has taught the skills to:
 - Distinguish functional programming and OOP.
 - Learn new programming languages quickly.
 - Master specific programming language concepts.
 - Evaluate programming languages and their constructs.
- Software is all about taking a few ideas (language constructs), and combining them with human ingenuity to create complex systems that run the world.